

## Cancelling a taskflow

Main idea- place a flag with each topology indicating if it is cancelled or not. The cancel method will be called by a modified return type of the `executor.run()` calls.

Details:

Create 2 atomic Boolean flags – `is_cancel` and `is_torn`. `is_cancel` will indicate if a topology has been cancelled. `is_torn` will indicate if a topology has been torn down.

Return type of executor run: Initially, executor runs returned `std::future` object. Now they will return `tf::Future` object. The new class will have a `std::future` member which will store the original future object. There will also be a topology in this class corresponding to the `executor.run()` call. The class will have overloading with the functions of `std::future` wherein the class will call the corresponding function of `std::future`.

e.g. `auto t1 = exec.run(tf1);`

`t1` will contain a `std::future<T>` `future_obj` and a `Topology`.

`t1.wait()` -> will call `future_obj.wait()`.

`t1.get()` -> will return `future_obj.get()`.

`T1.cancel()` -> will set `is_cancel` of the topology to 1.

For async tasks, there is no associated topology. So we create a `async_cancel` flag member in the `Node` class. Corresponding to async tasks, we have `t1.async_cancel()` to cancel async tasks.

## Cancelling a node:

Running of a node's task is done in `_invoke()` method of the `exec`. Corresponding to each task type, the corresponding invoke is done e.g. `_invoke_static_task`, `_invoke_dynamic_task`, etc. The checking of the flag is done in these functions. If the flag is 1, the task is not run.

e.g.

```
inline void Executor::_invoke_static_task(Worker& worker, Node* node) {
    _observer_prologue(worker, node);
    if (!node->_topology->is_cancel){
        std::get<Node::StaticTask>(node->_handle).work();
    }
    _observer_epilogue(worker, node);
}
```

Also, if a tpg is cancelled, the successors of the nodes will not be scheduled, the join\_counters will not be changed. Instead, the topology will be sent for tearing down directly.

Tearing a cancelled topology:

In conventional method, a topology will be rerun if the predicate is false, and the callback will be run after the topology run is complete. For a cancelled topology, there will be no checking of predicate and no callback function will be run. Instead, the topology will be torn down directly. We create a new function `_tear_cancelled_topology`. It is based on `_tear_down_topology` method. In this function, we first check if `is_torn` of the tpg is true or not. This flag indicates if that tpg has been torn down. If it hasn't been torn down, we acquire the lock and recheck if `is_torn` is false. Then we tear down the tpg just as `_tear_down_topology` does it. It makes `is_torn` true after tearing down to ensure that another thread does not tear down the same topology in parallel.

Overall, the following changes have to be done:

1. Executor.hpp-

Change return types of `run`, `run_n` and `run_until`;  
add flag checks to `invoke_static_task`, `invoke_dynamic_task`, etc;  
add a function `_tear_cancelled_topology`;  
add flag checks to call `_tear_cancelled_topology` in `_invoke` method;  
change return type of `async` method.

2. Taskflow.hpp

Add `tf::Future` class which has `std::future` object and topology. Add corresponding methods of `std::future` like `get()`, `wait()`, etc. Add `cancel()` method to set flag of topology.

3. Graph.hpp

Add `async_cancelled` flag to node class to check cancellation of async tasks.

4. Topology.hpp

Add `is_cancel` and `is_torn` flags to check if tpg is cancelled and torn down.